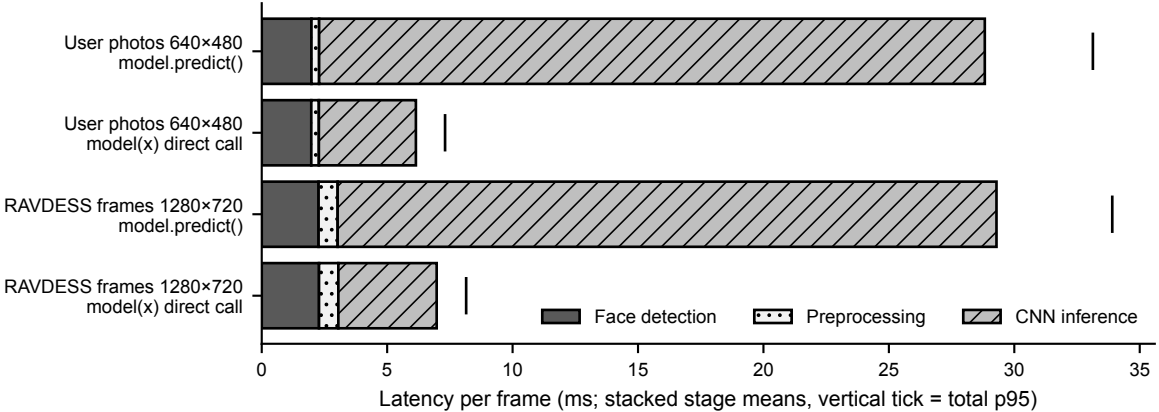
User photos 640×480
model.predict()
User photos 640×480
model(x) direct call
RAVDESS frames 1280×720
model.predict()
RAVDESS frames 1280×720
model(x) direct call
Face detection
Preprocessing
CNN inference
0
5
10
15
20
25
30
35
Latency per frame (ms; stacked stage means, vertical tick = total p95)